

STAT 576 Bayesian Analysis

Lecture 9: Hybrid Monte Carlo

Chencheng Cai

Washington State University

Hybrid Monte Carlo

- ▶ Hybrid Monte Carlo is a MCMC algorithm that combines the Metropolis algorithm with molecular dynamics (MD).
- ▶ The molecular dynamics (MD) is a **deterministic** algorithm that simulates the motion of particles in a physical system.
- ▶ We use the MD to propose a new state of the system, and then use the Metropolis algorithm to accept or reject the proposed state.
- ▶ Benefits:
 - ▶ It can explore the state space more efficiently than the random walk Metropolis.
 - ▶ It can handle high-dimensional state space.

Newton's Mechanics

Consider a single particle with mass m in a potential energy field $U(\mathbf{q})$.

- ▶ The particle's **position** is denoted by \mathbf{q} and its **momentum** by \mathbf{p} .
In **non-relativistic (NR)** mechanics, the momentum is given by

$$\mathbf{p} = m\mathbf{v},$$

where $\mathbf{v} = \dot{\mathbf{q}} := d\mathbf{q}/dt$ is the velocity.

- ▶ The **kinetic energy** of the particle is

$$T(\mathbf{p}) = \frac{\|\mathbf{p}\|^2}{2m} \stackrel{\text{NR}}{=} \frac{1}{2}m\|\mathbf{v}\|^2.$$

Newton's Mechanics

The (NR) dynamics of this particle is determined by Newton's second law:

$$\mathbf{F} = m\mathbf{a},$$

where \mathbf{F} is the force acting on the particle, and $\mathbf{a} = \ddot{\mathbf{q}} = d^2\mathbf{q}/dt^2$ is the acceleration.

- ▶ If the force is derived from a potential energy field $U(\mathbf{q})$, then the force is given by

$$\mathbf{F} = -\frac{\partial U(\mathbf{q})}{\partial \mathbf{q}}$$

- ▶ The $m\mathbf{a}$ is the time derivative of the momentum:

$$m\mathbf{a} = \dot{\mathbf{p}} = \frac{d\mathbf{p}}{dt}$$

- ▶ The Newton's second law in the **phase space** (\mathbf{q}, \mathbf{p}) is

$$-\frac{\partial U(\mathbf{q})}{\partial \mathbf{q}} = \dot{\mathbf{p}}.$$

Lagrangian Analytical Mechanics

We consider a (NR) particle in a system, whose dynamics is parametrized by $(t, \mathbf{q}, \dot{\mathbf{q}})$.

- We observe that

$$\frac{\partial T(\dot{\mathbf{q}})}{\partial \dot{\mathbf{q}}} = m\dot{\mathbf{q}}, \quad \mathbf{F} = -\frac{\partial U(\mathbf{q})}{\partial \mathbf{q}}$$

- By Newton's second law, we have

$$\frac{d}{dt} \frac{\partial T(\dot{\mathbf{q}})}{\partial \dot{\mathbf{q}}} + \frac{\partial U(\mathbf{q})}{\partial \mathbf{q}} = 0$$

- If we define the **Lagrangian** of the system by

$$L(t, \mathbf{q}, \dot{\mathbf{q}}) = T(\dot{\mathbf{q}}) - U(\mathbf{q}),$$

then the equation above is equivalent to the **Euler-Lagrange equation**:

$$\frac{\partial L}{\partial \mathbf{q}} - \frac{d}{dt} \frac{\partial L}{\partial \dot{\mathbf{q}}} = 0$$

Lagrangian Analytical Mechanics

We consider a particle in a system, whose dynamics is parametrized by $(t, \mathbf{q}, \mathbf{p})$.

- By Newton's second law, we have

$$-\frac{\partial U(\mathbf{q})}{\partial \mathbf{q}} = \dot{\mathbf{p}}.$$

- On the other hand, we have

$$\frac{\partial T(\mathbf{p})}{\partial \mathbf{p}} = \dot{\mathbf{q}}$$

- Combining the two equations above, we have the **Hamilton's equations**:

$$\begin{aligned}\dot{\mathbf{q}} &= \frac{\partial H(\mathbf{q}, \mathbf{p})}{\partial \mathbf{p}}, \\ \dot{\mathbf{p}} &= -\frac{\partial H(\mathbf{q}, \mathbf{p})}{\partial \mathbf{q}},\end{aligned}$$

where $H(\mathbf{q}, \mathbf{p}) = T(\mathbf{p}) + U(\mathbf{q})$ is the **Hamiltonian** of the system.

Lagrangian Analytical Mechanics

The Hamiltonian's equations can be derived from the Euler-Lagrange equation directly.

- ▶ We define the momentum \mathbf{p} as the conjugate variable of the position \mathbf{q} :

$$\mathbf{p} = \frac{\partial L}{\partial \dot{\mathbf{q}}}$$

- ▶ The Hamiltonian is defined as the Legendre transform of the Lagrangian:

$$H(\mathbf{q}, \mathbf{p}) = \mathbf{p} \cdot \dot{\mathbf{q}} - L(\mathbf{q}, \dot{\mathbf{q}})$$

- ▶ The Hamiltonian's equations are derived from the Euler-Lagrange equation:

$$\begin{aligned}\frac{\partial H}{\partial \mathbf{p}} &= \dot{\mathbf{q}}, \\ \frac{\partial H}{\partial \mathbf{q}} &= -\frac{\partial L}{\partial \mathbf{q}} = -\frac{d}{dt} \frac{\partial L}{\partial \dot{\mathbf{q}}} = -\dot{\mathbf{p}}.\end{aligned}$$

Molecular Dynamics Simulation

Consider simulating a system with Hamiltonian $H(\mathbf{q}, \mathbf{p})$.

- Use Taylor expansion:

$$\begin{aligned}\mathbf{q}(t + dt) &= \mathbf{q}(t) + \dot{\mathbf{q}}(t)dt + \frac{1}{2}\ddot{\mathbf{q}}(t)(dt)^2 + \cdots, \\ &= \mathbf{q}(t) + \frac{\mathbf{p}(t)}{m}dt + \frac{1}{2}\frac{\dot{\mathbf{p}}(t)}{m}(dt)^2 + \cdots, \\ \mathbf{p}(t + dt) &= \mathbf{p}(t) + \dot{\mathbf{p}}(t)dt + \frac{1}{2}\ddot{\mathbf{p}}(t)(dt)^2 + \cdots.\end{aligned}$$

- Furthermore, we have

$$\begin{aligned}\mathbf{q}(t + dt) + \mathbf{q}(t - dt) &= 2\mathbf{q}(t) + \frac{\dot{\mathbf{p}}(t)}{m}(dt)^2 + O((dt)^4), \\ \mathbf{q}(t + dt) - \mathbf{q}(t - dt) &= 2\mathbf{p}(t)dt + O((dt)^3).\end{aligned}$$

Molecular Dynamics Simulation

Størmer-Verlet algorithm:

- ▶ At time t with a time increment Δt .
- ▶ We update the position by:

$$\mathbf{q}(t + \Delta t) = 2\mathbf{q}(t) - \mathbf{q}(t - \Delta t) - \frac{(\Delta t)^2}{m} \left. \frac{\partial H}{\partial \mathbf{q}} \right|_t$$

- ▶ We update the momentum by:

$$\mathbf{p}(t + \Delta t) = m \frac{\mathbf{q}(t + \Delta t) - \mathbf{q}(t - \Delta t)}{2\Delta t}$$

or

$$\mathbf{p}(t + \Delta t) = m \frac{\mathbf{q}(t + \Delta t) - \mathbf{q}(t)}{\Delta t}$$

Molecular Dynamics Simulation

Leap-frog algorithm:

- ▶ At time t with a time increment Δt .
- ▶ We update the position by:

$$\mathbf{q}(t + \Delta t) = \mathbf{q}(t) + \Delta t \frac{\mathbf{p}\left(t + \frac{1}{2}\Delta t\right)}{m}$$

- ▶ We update the momentum at the half-time interval by:

$$\mathbf{p}\left(t + \frac{1}{2}\Delta t\right) = \mathbf{p}\left(t - \frac{1}{2}\Delta t\right) - \left.\frac{\partial H}{\partial \mathbf{q}}\right|_t \Delta t$$

Molecular Dynamics Simulation

Leap-frog algorithm (alternative form):

- ▶ At time t with a time increment Δt .
- ▶ We update the momentum by the first half-time interval:

$$\mathbf{p}\left(t + \frac{1}{2}\Delta t\right) = \mathbf{p}(t) - \left.\frac{\partial H}{\partial \mathbf{q}}\right|_t \frac{\Delta t}{2}$$

- ▶ We update the position by:

$$\mathbf{q}(t + \Delta t) = \mathbf{q}(t) + \Delta t \frac{\mathbf{p}\left(t + \frac{1}{2}\Delta t\right)}{m}$$

- ▶ We update the momentum by the second half-time interval:

$$\mathbf{p}(t + \Delta t) = \mathbf{p}\left(t + \frac{1}{2}\Delta t\right) - \left.\frac{\partial H}{\partial \mathbf{q}}\right|_{t+\Delta} \frac{\Delta t}{2}$$

Molecular Dynamics Simulation

Størmer-Verlet algorithm and Leap-frog algorithm are identical algorithms.

- From the position update set in leap-frog algorithm, we have

$$\mathbf{q}(t + \Delta t) = \mathbf{q}(t) + \Delta t \frac{\mathbf{p}(t + \frac{1}{2}\Delta t)}{m}$$
$$\mathbf{q}(t) = \mathbf{q}(t - \Delta t) + \Delta t \frac{\mathbf{p}(t - \frac{1}{2}\Delta t)}{m}$$

- Substitute the second equation into the first equation, we have

$$\mathbf{q}(t + \Delta t) = 2\mathbf{q}(t) - \mathbf{q}(t - \Delta t) - \frac{(\Delta t)^2}{m} \left. \frac{\partial H}{\partial \mathbf{q}} \right|_t$$

which is the same as the position update in Størmer-Verlet algorithm.

Molecular Dynamics Simulation

The Hamiltonian dynamics has the following preservation properties:

- ▶ The Hamiltonian $H(\mathbf{q}, \mathbf{p})$ is preserved by the dynamics.

$$\frac{dH}{dt} = \frac{\partial H}{\partial \mathbf{q}} \dot{\mathbf{q}} + \frac{\partial H}{\partial \mathbf{p}} \dot{\mathbf{p}} = 0.$$

- ▶ The volume in the phase space is preserved by the dynamics.

Let $V(t) = \{(\mathbf{q}(t), \mathbf{p}(t)) : (\mathbf{q}(0), \mathbf{p}(0)) \in V(0)\}$, then

Louville's theorem: $V(t)$ has the same volume as $V(0)$.

The MD simulation is volume-preservation and time reversible, but not Hamiltonian-preserving.

- ▶ The MD simulation is volume-preserving because the Jacobian of the transformation at each step is 1.
- ▶ The MD simulation is time reversible because we can simulate the system backward by reversing the momentum.
- ▶ Hamiltonian is not preserved because of the discretization error.

Example

Consider a dynamic system with

$$U(q) = q^2 - 2 \log[\cosh(2q)] + 3, \quad T(p) = p^2/2$$

The Hamiltonian is


$$H(q, p) = U(q) + T(p) = q^2 + \frac{p^2}{2} - 2 \log[\cosh(2q)]$$

The Hamiltonian's equations give

$$\dot{p} = -\frac{\partial H}{\partial q} = -2q + 4 \tanh(2q)$$

We simulate the system using the leap-frog algorithm with $\Delta t = 0.1$, starting position $q(0) = 1$, and five different momenta $p(0) = 1, 2, 3, 4, 5$.

Example



Example

```
n = 200
dt = 0.1


phase_sample = array(dim=c(5, 2, n+1))
phase_sample[,1,1] = 1
phase_sample[,2,1] = 1:5

pdot = function(q) {-2*q + 4*tanh(2*q)}

for(t in 1:n){
  p = phase_sample[,2,t] + 0.5 * pdot(phase_sample[,1,t]) * dt
  phase_sample[,1,t+1] = phase_sample[,1,t] + dt * p
  phase_sample[,2,t+1] = p + 0.5 * pdot(phase_sample[,1,t+1]) * dt
}
```


Example

Trajectories in the phase space of the system with different initial momentums.



Hamiltonian Monte Carlo

Suppose we have a target distribution $\pi(\mathbf{x}) \propto \exp\{-U(\mathbf{x})\}$, and we want to sample from it.

- ▶ We introduce an auxiliary momentum variable \mathbf{p} with density $p(\mathbf{p}) = \exp\{-T(\mathbf{p})\}$.
- ▶ The joint distribution of (\mathbf{x}, \mathbf{p}) is

$$p(\mathbf{x}, \mathbf{p}) = \exp\{-H(\mathbf{x}, \mathbf{p})\},$$

where $H(\mathbf{x}, \mathbf{p}) = U(\mathbf{x}) + T(\mathbf{p})$ is the Hamiltonian.

Observations:

- ▶ If we can sample from the joint distribution $p(\mathbf{x}, \mathbf{p})$, then we can sample from the target distribution $\pi(\mathbf{x})$ by marginalizing out \mathbf{p} .
- ▶ The Hamiltonian dynamics is reversible, meaning that if $(\mathbf{x}(t), \mathbf{p}(t))$ results in $(\mathbf{x}(t + \Delta t), \mathbf{p}(t + \Delta t))$, then $(\mathbf{x}(t + \Delta t), -\mathbf{p}(t + \Delta t))$ results in $(\mathbf{x}(t), -\mathbf{p}(t))$.

Hamiltonian Monte Carlo

Suppose we have a target distribution $\pi(\mathbf{x}) \propto \exp\{-U(\mathbf{x})\}$. The Hamiltonian Monte Carlo algorithm is as follows: (for time $t + 1$)

- ▶ Sample a momentum \mathbf{p} from

$$p(\mathbf{p}) \propto \exp\{-T(\mathbf{p})\}$$

- ▶ Run the leap-frog algorithm for L steps with step size ϵ jumping from (\mathbf{x}, \mathbf{p}) to $(\mathbf{x}', \mathbf{p}')$.
- ▶ Accept the new state $(\mathbf{x}', \mathbf{p}')$ with probability

$$\alpha = \min \{1, \exp\{-H(\mathbf{x}', \mathbf{p}') + H(\mathbf{x}, \mathbf{p})\}\}$$

Remark: The leap-frog step can be replaced by any deterministic time-reversible and volume-preserving dynamics.

Hamiltonian Monte Carlo

Justification for correctness:

- ▶ (time-reversible) If the L leap-frog steps map (\mathbf{x}, \mathbf{p}) to $(\mathbf{x}', \mathbf{p}')$, then they also map $(\mathbf{x}', -\mathbf{p}')$ to $(\mathbf{x}, -\mathbf{p})$.
- ▶ (volume-preserving) The leap-frog algorithm is volume-preserving. Therefore $d\mathbf{x}d\mathbf{p} = d\mathbf{x}'d\mathbf{p}'$.
- ▶ The acceptance probability in a Metropolis-Hastings algorithm is

$$\alpha = \min \left\{ 1, \frac{\pi(\mathbf{x}')}{\pi(\mathbf{x})} \frac{p(\mathbf{x}', \mathbf{x})d\mathbf{x}'d\mathbf{p}'}{p(\mathbf{x}, \mathbf{x}')d\mathbf{x}d\mathbf{p}} \right\} = \min \left\{ 1, \frac{\pi(\mathbf{x}')}{\pi(\mathbf{x})} \frac{p(\mathbf{p}')}{p(\mathbf{p})} \right\} = \min \left\{ 1, \frac{e^{-H(\mathbf{x}', \mathbf{p})}}{e^{-H(\mathbf{x}, \mathbf{p})}} \right\}$$

Example

Consider sampling from the $\text{Beta}(2, 2)$ distribution. The Hamiltonian is

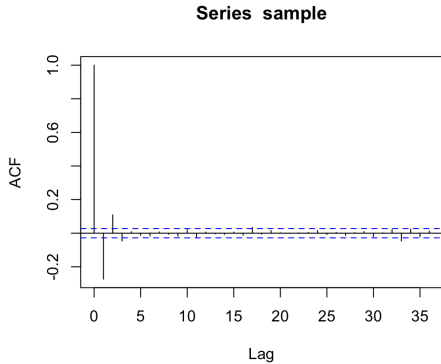

$$H(x, p) = -\log(x(1-x)) + \frac{p^2}{2}$$

```
pdot = function(x){1/x-1/(1-x)}  
h = function(x, p){  
  if(x*(1-x) <= 0) return(Inf)  
  else return(-log(x*(1-x)) + p**2/2)}  
  
eps = 0.05  
L = 10  
n = 5000  
burnin = 500  
sample = rep(0, n+1)  
x = 0.5
```

Example

```
for(i in 1:(n+burnin)){  
  p = rnorm(1)  
  h_old = h(x, p)  
  y = x  
  for(j in 1:L){  
    p = p + 0.5 * pdot(y) * eps  
    y = y + p * eps  
    p = p + 0.5 * pdot(y) * eps  
  }  
  h_new = h(y, p)  
  if(runif(1) <= exp(-h_new + h_old)) x = y  
  if(i > burnin) sample[i-burnin] = x  
}
```

Example



Langevin Dynamics

Recall the Taylor's expansion for the position:

$$\mathbf{q}(t + dt) = \mathbf{q}(t) + \frac{\mathbf{p}(t)}{m}dt + \frac{1}{2} \frac{\dot{\mathbf{p}}(t)}{m}(dt)^2 + \dots$$

- ▶ With Hamiltonian's equations, we have

$$\mathbf{q}(t + dt) = \mathbf{q}(t) + \frac{\mathbf{p}(t)}{m}dt - \frac{1}{2} \frac{\partial H / \partial \mathbf{q}}{m}(dt)^2 + \dots$$

- ▶ In the Hamiltonian Monte Carlo, we use the leap-frog algorithm to simulate the dynamics:
 - ▶ $m = 1$
 - ▶ \mathbf{p} is drawn from a standard multivariate normal.
 - ▶ $H = U(\mathbf{q}) + \|\mathbf{p}\|^2$.

Langevin Dynamics

If the target distribution is $\mathbf{x} \sim \pi(\mathbf{x})$, one step of leap-frog is equivalent to the following Langevin dynamics:

$$\mathbf{x}_{t+1} = \mathbf{x}_t + \frac{1}{2} \frac{\partial \log \pi(\mathbf{x})}{\partial \mathbf{x}} h + \sqrt{h} \mathbf{Z}_t,$$

- ▶ Time indices are relabelled to integers.
- ▶ \mathbf{Z}_t is standard multivariate normal. (random sampling for \mathbf{p})
- ▶ h is the step size. ($\sqrt{dt} = h$).

The corresponding stochastic differential equation is

$$d\mathbf{x}_t = \frac{1}{2} \frac{\partial \log \pi(\mathbf{x})}{\partial \mathbf{x}} dt + d\mathbf{W}_t.$$

Langevin Dynamics

The **Langevin dynamics** is a continuous-time stochastic process for a particle in a potential field $U(\mathbf{x})$. The stochastic differential equation is

$$d\mathbf{x}_t = -\frac{1}{2} \frac{\partial U(\mathbf{x})}{\partial \mathbf{x}} dt + d\mathbf{W}_t,$$

where \mathbf{W}_t is standard Brownian motion.

Let $p(\mathbf{x}, t)$ be the density of the particle at time t .

The **Fokker-Planck equation** (also known as the **Kolmogorov forward equation**) is

$$\frac{\partial p(\mathbf{x}, t)}{\partial t} = \frac{1}{2} \nabla \cdot [p(\mathbf{x}, t) \nabla U(\mathbf{x}) + \nabla p(\mathbf{x}, t)].$$

$p(\mathbf{x}, t) = \exp\{-U(\mathbf{x})\}$ is the stationary distribution such that

$$\frac{\partial p(\mathbf{x}, t)}{\partial t} = 0$$

Langevin Dynamics

To generate samples from the target distribution $\pi(\boldsymbol{x})$, we can simulate the Langevin dynamics with SDE:

$$d\boldsymbol{x}_t = \frac{1}{2} \nabla \log \pi(\boldsymbol{x}) dt + d\boldsymbol{W}_t.$$

Its discretized version is

$$\boldsymbol{x}_{t+1} = \boldsymbol{x}_t + \frac{1}{2} \nabla \log \pi(\boldsymbol{x}_t) h + \sqrt{h} \boldsymbol{Z}_t,$$

where \boldsymbol{Z}_t is standard multivariate normal.

This step (without the auxiliary momentum) can replace the leap-frog step in the Hamiltonian Monte Carlo.

Metropolis-adjusted Langevin Algorithm

Suppose the target distribution is $\pi(\mathbf{x}) \propto \exp\{-U(\mathbf{x})\}$. The Metropolis-adjusted Langevin algorithm is as follows: (for time $t + 1$)

- ▶ Run the Langevin dynamics with step size ϵ jumping from \mathbf{x} to \mathbf{x}' .

$$\mathbf{x}' = \mathbf{x} - \frac{1}{2} \nabla U(\mathbf{x}) h + \sqrt{h} \mathbf{Z}_t.$$

- ▶ Accept the new state \mathbf{x}' with probability

$$\alpha = \min \left\{ 1, \frac{\pi(\mathbf{x}') \mathcal{N}(\mathbf{x} \mid \mathbf{x}' - \frac{h}{2} \nabla U(\mathbf{x}'), h \mathbf{I})}{\pi(\mathbf{x}) \mathcal{N}(\mathbf{x}' \mid \mathbf{x} - \frac{h}{2} \nabla U(\mathbf{x}), h \mathbf{I})} \right\}$$

Example



Draw samples from the $\text{Beta}(2, 2)$ distribution using the Langevin dynamics.
We have

$$U(x) = -\log x(1-x) \text{ and } dU/dx = -\frac{1}{x} + \frac{1}{1-x}$$

Example

```
n = 10000
burnin = 1000
sample = rep(0, n)
x = 0.5
h = 0.1
hsqrt = sqrt(h)
for(i in 1:(n+burnin)){
  z = rnorm(1)
  y = x - 0.5*(2*x-1)/x/(1-x)*h + hsqrt*z
  if(y*(1-y)>0){
    logpy = log(y*(1-y)) - 0.5*(x-y+0.5*(2*y-1)/y/(1-y)*h)**2/h
    logpx = log(x*(1-x)) - 0.5*z**2
    if(runif(1) <= exp(logpy-logpx)) x = y
  }
  if(i > burnin) sample[i - burnin] = x
}
```

Example



Langevin Dynamics

The Langevin dynamics can also be written as the following stochastic differential equation:

$$d\mathbf{x}_t = -\nabla U(\mathbf{x})dt + \sqrt{2}d\mathbf{W}_t,$$

where \mathbf{W}_t is standard Brownian motion.

With the discretization as

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \nabla U(\mathbf{x}_t)h + \sqrt{2h}\mathbf{Z}_t.$$

This version (along with the previous version) is also called the **overdamped Langevin dynamics**.

Langevin Dynamics

The **underdamped Langevin dynamics** is a second-order stochastic differential equation:

$$d\mathbf{x}_t = \mathbf{v}_t dt, \quad d\mathbf{v}_t = -\nabla U(\mathbf{x}_t) dt - \gamma \mathbf{v}_t dt + \sqrt{2\gamma} d\mathbf{W}_t,$$

where γ is the friction coefficient, and the mass is assumed to be 1.

- ▶ When $\gamma \gg 1$, the dynamics is overdamped. When $\gamma \ll 1$, the dynamics is underdamped.
- ▶ The overdamped Langevin dynamics converges slower than the underdamped Langevin dynamics.

Example: the Eight Schools

Recall the eight schools example. The hierarchical model is

$$p(\mu, \tau) \propto 1$$

$$\alpha_j \sim \mathcal{N}(\mu, \tau^2)$$

$$y_j \sim \mathcal{N}(\alpha_j, \sigma_j^2)$$

The posterior distribution is $\theta = (\mu, \tau, \alpha)$.

$$p(\theta \mid y) \propto p(\mu, \tau) \prod_{j=1}^8 p(\alpha_j \mid \mu, \tau) p(y_j \mid \alpha_j, \sigma_j) \propto \frac{1}{\tau^J} e^{-\frac{\sum (\alpha_j - \mu)^2}{2\tau^2}} e^{-\sum \frac{(\alpha_j - y_j)^2}{2\sigma_j^2}}$$

Example: the Eight Schools

In both HMC and LMC, we need to compute the gradient of the log posterior.

$$\frac{\partial \log p(\theta \mid y)}{\partial \mu} = -\frac{J}{\tau^2}(\mu - \bar{\alpha})$$

$$\frac{\partial \log p(\theta \mid y)}{\partial \tau} = -\frac{J}{\tau} + \frac{1}{\tau^3} \sum_{j=1}^J (\alpha_j - \mu)^2$$

$$\frac{\partial \log p(\theta \mid y)}{\partial \alpha_j} = -\frac{\alpha_j - \mu}{\tau^2} - \frac{\alpha_j - y_j}{\sigma_j^2}$$

The Hamiltonian is

$$H(\mu, \tau, \boldsymbol{\alpha}, p_\mu, p_\tau, \mathbf{p}_\alpha) = -\log p(\theta \mid y) + \frac{1}{2\sigma_p^2} (p_\mu^2 + p_\tau^2 + \|\mathbf{p}_\alpha\|^2)$$

Example: the Eight Schools

```
n = 1000
burnin = 100
J = 8

mu.sample = rep(0, n)
tau.sample = rep(0, n)
alpha.sample = array(dim=c(J, n))

mu = 0
tau = 5
alpha = rep(0, J)

L = 20
eps = 0.05
p.sigma = 0.1
```

Example: the Eight Schools

```
ham = function(mu, tau, alpha, p.mu, p.tau, p.alpha){  
  if(tau <= 0)  
    return(Inf)  
  p1 = 0.5*sum((alpha-mu)**2)/tau**2  
  p2 = 0.5*sum((alpha-y)**2/s**2)  
  p3 = J*log(tau)  
  p4 = p.mu**2/2/p.sigma**2  
  p5 = p.tau**2/2/p.sigma**2  
  p6 = sum(p.alpha**2)/2/p.sigma**2  
  return(p1+p2+p3+p4+p5+p6)  
}
```

Example: the Eight Schools

```
for(i in 1:(n+burnin)){
  p.mu = rnorm(1) * p.sigma
  p.tau = rnorm(1) * p.sigma
  p.alpha = rnorm(J) * p.sigma
  h_old = ham(mu, tau, alpha, p.mu, p.tau, p.alpha)
  n.mu = mu
  n.tau = tau
  n.alpha = alpha
  // leap-frog step here
  h_new = ham(n.mu, n.tau, n.alpha, p.mu, p.tau, p.alpha)
  if(runif(1) <= exp(h_old - h_new)){
    mu = n.mu
    tau = n.tau
    alpha = n.alpha
  }
  if(i > burnin){
    mu.sample[i-burnin] = mu
    tau.sample[i-burnin] = tau
    alpha.sample[:,i-burnin] = alpha
  }
}
```

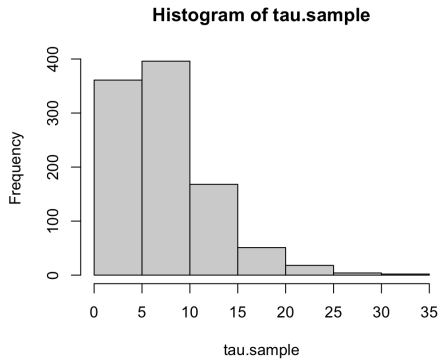

Example: the Eight Schools

```
// leap-frog step
for(j in 1:L){
  p.mu = p.mu + 0.5*eps* (-J) * (n.mu - mean(n.alpha)) / n.tau ** 2
  p.tau = p.tau+0.5*eps* (-J/n.tau+sum((n.alpha-n.mu)**2)/n.tau**3)
  p.alpha = p.alpha+0.5*eps*(-(n.alpha-n.mu)/n.tau**2-(n.alpha-y)/s**2)

  n.mu = n.mu + p.mu * eps / p.sigma ** 2
  n.tau = n.tau + p.tau * eps / p.sigma ** 2
  n.alpha = n.alpha + p.alpha * eps / p.sigma ** 2


  p.mu = p.mu + 0.5*eps* (-J) * (n.mu - mean(n.alpha)) / n.tau ** 2
  p.tau = p.tau+0.5*eps* (-J/n.tau+sum((n.alpha-n.mu)**2)/n.tau**3)
  p.alpha = p.alpha+0.5*eps*(-(n.alpha-n.mu)/n.tau**2-(n.alpha-y)/s**2)
}
```

Example: the Eight Schools




Example: the Eight Schools


Histogram of alpha.sample[1,]




Histogram of alpha.sample[2,]




Histogram of alpha.sample[3,]




Histogram of alpha.sample[4,]




Histogram of alpha.sample[5,]




Histogram of alpha.sample[6,]



Histogram of alpha.sample[7,]



Histogram of alpha.sample[8,]



Example: the Eight Schools

We can also use the Langevin dynamics to sample from the posterior distribution.

```
h = 0.1
hsqrt = sqrt(h)
mu = 0
tau = 5
alpha = rep(0, J)

logp = function(mu, tau, alpha){
  if(tau <= 0) return(-Inf)
  p1 = 0.5*sum((alpha-mu)**2)/tau**2
  p2 = 0.5*sum((alpha-y)**2/s**2)
  p3 = J*log(tau)
  return(-(p1+p2+p3))
}

fun.dlogp = function(mu, tau, alpha){
  dmu = -J*(mu - mean(alpha))/tau ** 2
  dtau = -J/tau+sum((alpha-mu)**2)/tau**3
  dalpha = -(alpha - mu)/tau**2 - (alpha-y)/s**2
  return(list(mu=dmu, tau=dtau, alpha=dalpha))
}
```

Example: the Eight Schools

```
for(i in 1:(n+burnin)){
  z.mu = rnorm(1)
  z.tau = rnorm(1)
  z.alpha = rnorm(J)
  dlogp = fun.dlogp(mu, tau, alpha)
  n.mu = mu + 0.5*h*dlogp$mu + hsqrt * z.mu
  n.tau = tau + 0.5*h*dlogp$tau + hsqrt*z.tau
  n.alpha = alpha + 0.5*h*dlogp$alpha + hsqrt*z.alpha
  dlogp = fun.dlogp(n.mu, n.tau, n.alpha)
  loga = logp(n.mu, n.tau, n.alpha) - logp(mu, tau, alpha)
  loga = loga - 0.5 * (mu - n.mu - 0.5*h*dlogp$mu)**2/h
  loga = loga - 0.5 * (tau - n.tau - 0.5*h*dlogp$tau)**2/h
  loga = loga - 0.5 * sum((alpha - n.alpha - 0.5*h*dlogp$alpha)**2/h)
  loga = loga + 0.5 * z.mu ** 2 + 0.5*z.tau**2 + 0.5 * sum(z.alpha**2)

  //Accept with prob exp(loga)
  //store the sample
}
```